

How to implement mutex?

```
1  class LockOne implements Lock {  
2      private boolean[] flag = new boolean[2];  
3  
4      public void lock() {  
5          int i = ThreadID.get();  
6          int j = 1 - i;  
7          flag[i] = true;  
8          while (flag[j]) {} // wait  
9      },  
10  
11     public void unlock() {  
12         int i = ThreadID.get();  
13         flag[i] = false;  
14     }  
15 }
```

How to implement mutex? Option 2

```
1 class LockTwo implements Lock {  
2     private volatile int victim;  
3  
4     public void lock() {  
5         int i = ThreadID.get();  
6         victim = i;  
7  
8         while(victim == i) {} // wait  
9     }  
10  
11     public void unlock() {}  
12 }
```

Working version

```
1 class Peterson implements Lock {
2     private volatile boolean[] flag = new boolean[2];
3     private volatile int victim;
4
5     public void lock() {
6         int i = ThreadID.get();
7         int j = 1 - i;
8         flag[i] = true;
9         victim = i;
10        while (flag[j] && victim == i) {}; // wait
11    }
12
13    public void unlock() {
14        int i = ThreadID.get();
15        flag[i] = false;
16    }
17 }
```

Synchronization primitives

Non-recursive - when re-capturing by the same thread, call **deadlock**

Recursive - allow you to capture yourself by the same thread multiple times

Recursive primitives

Lock

Lock

Lock

Unlock

Unlock

Unlock

Method1: Lock

Method2: Lock

Method3: Lock

Method4: Unlock

Method5: Lock

Method6: Unlock

Method7: Unlock

Method8: Unlock

Example 1

```
1  Class Vector {  
2      Mutex m;  
3  
4      public void add() {  
5          m.lock();  
6          size();  
7          extend();  
8          m.unlock();  
9      }  
10  
11     public int size() {  
12         m.lock();  
13         int size = getSize();  
14         m.unlock()  
15  
16         return size;  
17     }  
18 }
```

Example 2

```
1 | Class Vector {  
2 |     Mutex m;  
3 |  
4 |     public void add() {  
5 |         m.lock();  
6 |         unsafeAdd();  
7 |         m.unlock();  
8 |     }  
9 |  
10 |    public int size() {  
11 |        m.lock();  
12 |        size = unsafeSize()  
13 |        m.unlock()  
14 |  
15 |        return size;  
16 |    }  
17 |  
18 |    public void unsafeAdd() {  
19 |        unsafeSize();  
20 |        extend();  
21 |    }  
22 |  
23 |    public int unsafeSize() {  
24 |        return getSize();  
25 |    }  
26 | }
```

Types of mutexes

- Timed mutex
- Shared mutex
- Spin mutex
- Futex

Timed mutex

Tries to capture the mutex within the specified time

Shared mutex

Allows you to **lock** read-only, write-only, or mixed. Allows you to collapse the **lock-on** separate operations into a single lock.

w r r r w r r w r r r = w r w r w r

Spin mutex

Active waiting

java.util.concurrent.atomic.AtomicInteger

```
159     public final int ↴ getAndIncrement() {  
160         for (;;) {  
161             int current = get();  
162             int next = current + 1;  
163             if (compareAndSet(current, next))  
164                 return current;  
165         }  
166     }
```

CAS - operations

CAS - compare and set, compare and swap

```
bool compare_and_set( int *source, int oldValue, int newValue)
```

- Returns whether the value was set successfully
- Atomic at the processor level (i486+): **cmpxchg**

Examples:

```
a = 5; current_value = a;
```

```
compare_and_set(&a, current_value, 6)
```

SPIN lock - Active waiting

Use a lot of CPU time

When to use?

SPIN lock - Active waiting

Spinlock can be better when you plan to hold the lock for an extremely short interval (for example to do nothing but increment a counter), and contention is expected to be rare.

Benefits:

- On unlock, there is no need to check if other threads may be waiting for the lock and waking them up. Unlocking is simply a single atomic write instruction.
- Failure to immediately obtain the lock does not put your thread to sleep, so it may be able to obtain the lock with much lower latency as soon as it does become available.