

Threads

Thread

Threading is the creation and management of a variety of executable elements within a single process

A process contains one or more threads

Exist in the same process address space

The advantages of threads

- Parallelism
- Context switch
- Memory saving

What to choose?

Processes are generally **more reliable**

- processes interacting at different times
- distributed application
- data security

Threads - generally **faster**

- easier IPC

Types of implementations

- User threads (Linux 2.4)
- Kernel thread (Linux version ≥ 2.6 , NPTL)
- Hybrid implementation

User threads

The benefits of custom threads are as follows:

- thread switching does not require kernel participation - no switching from task mode to kernel mode;
- planning can be determined by the application-the best algorithm is chosen;
- user threads can be used on any OS - only a compatible thread library is required.

Disadvantages of user threads:

- most system calls are blocking and the kernel blocks processes-including all threads within the process;
- the kernel can only schedule processes to processors - two threads within the same process cannot run simultaneously on two different processors.

Kernel thread

Benefits of kernel-level threads:

- the kernel can schedule multiple threads of the same process to run on multiple processors at the same time, blocking is performed at the thread level;
- kernel procedures can be multithreaded.

Disadvantages:

- switching threads within a single process requires the participation of the kernel.

POSIX Threads

There are many implementations of threads:

Android, Apache, GNOME, Mozilla

Unified standard of UNIX-threads

POSIX-Threads

API

`<pthread.h>`

About 100 system calls:


- thread control-functions for creating, destroying, joining, and detaching threads
- synchronization - functions to control thread synchronization

Linking

Thread library: **libpthread**

gcc -pthread example.c

Thread life cycle



libc: pthread_cancel.c

```
107     /* Mark the thread as canceled. This has to be done
108        atomically since other bits could be modified as well. */
109     while (atomic_compare_and_exchange_bool_acq (&pd->cancelhandling, newval,
110                                                  oldval));
111
112     return result;
113 }
```

https://github.com/lattera/glibc/blob/a2f34833b1042d5d8eeb263b4cf4caaea138c4ad/nptl/pthread_cancel.c

Creating threads

```
#include <pthread.h>
```

```
int pthread_create (pthread_t *thread,  
                   const pthread_attr_t *attr,  
                   void *(*start_routine) (void *),  
                   void *arg);
```

Return result

0 - on success

otherwise, an error number (without the use of errno!)

EAGAIN	the calling process is severely under-resourced to create a new thread; this is usually because the process has reached the thread count limit for each user or for the entire system;
EINVAL	the pthread_attr_t object, specified via the attr has an invalid attributes;
EPERM	the calling process does not have permission to set some attributes of the pthread_attr_t object specified through attr

Example

```
pthread_t thread;  
int ret;  
ret = pthread_create (&thread, NULL, start_routine, NULL);  
if (!ret) {  
    errno = ret;  
    perror("pthread_create");  
    return -1;  
}
```

Thread ID

Unlike **PID-a** is assigned by the library

Although there is **gettid()** - use it does not make sense

```
#include <pthread.h>
```

```
pthread_t pthread_self (void);
```


Comparison of identifiers

pthread_t - thread ID

```
#include <pthread.h>
```

```
int pthread_equal (pthread_t t1, pthread_t t2);
```

0 - if equal, otherwise not equal

The end of the stream

- if a thread returns from the start procedure, it is interrupted; this is the equivalent of "out of bounds" in main();
- if the thread invokes the function **pthread_exit()**, it terminates; this is the equivalent of calling exit();
- if a thread is canceled by another thread through the **pthread_cancel()** function, it terminates; this is analogous to sending a SIGKILL signal through kill().

Exit

```
#include <pthread.h>
```

```
void pthread_exit (void *retval);
```

Join threads

```
#include <pthread.h>
```

```
int pthread_join (pthread_t thread, void **retval);
```

Some analog of call **wait()**

retval - returned value

Each thread can join any other

Possible errors

EDEADLK — a deadlock has occurred — thread is already waiting to join the calling thread or is itself a calling thread;

EINVAL — it is impossible to attach a thread, specified through the **thread**

ESRCH — the value of the **thread** is unacceptable.

Example

Create threads, attach

pthread1.c

pthread1-2.c

pthread1-3.c

pthread1-4.c

Exercise 4a

Create program that create two threads.

First thread should output: "Hello"

Second thread should output: "World"

Main thread should wait (*use `pthread_join`*) while all threads are finished

templates: pthread1-*.c

Example

PI computation without race conditions

pthread_pi_race.c

pthread_pi.c

Initializing mutexes

```
#include <pthread.h>
```

```
pthread_mutex_t mutex =  
PTHREAD_MUTEX_INITIALIZER;
```

Locking mutexes

```
#include <pthread.h>
```

```
int pthread_mutex_lock (pthread_mutex_t *mutex);
```

A successful call will block the calling thread until the mutex specified as mutex is available

Errors

EDEADLK — the calling thread already owns the requested mutex

EINVAL — mutex value is not valid.

Unlocking mutexes

```
#include <pthread.h>
```

```
int pthread_mutex_unlock (pthread_mutex_t *mutex);
```

A successful call releases the mutex specified by mutex and return 0

Errors

EINVAL — mutex value is not valid

EPERM — the calling process does not own the mutex specified as mutex; attempting to release a mutex that you do not own is an error.

Example

PI computation with single shared memory and race conditions

pthread_pi_lock_1.c - race

pthread_pi_lock_2.c - lock

pthread_pi_lock_3.c - smart lock

pthread_pi_lock_4.c - spin lock

pthread_pi_lock_5.c - spin lock 2

Semaphores

// Init

```
int sem_init(sem_t *sem, int pshared, unsigned int value);
```

// Increase by 1

```
int sem_post(sem_t *sem);
```

// Decrease by 1

```
int sem_wait(sem_t *sem);
```

Example

Basic semaphore usage

`pthread_sems.c`

Exercise 4b

Create parallel version of program for vectors multiplication based on PThreads.

Questions:

1. Compare execution time of the two versions of the program:
 - a. based on **forks**
 - b. based on **pthread**s.

Template: vector-multiplication-fork.c

Condition Variables

Condition variables provide yet another way for threads to synchronize.

While mutexes implement synchronization by controlling **thread access to data**, condition variables allow threads to synchronize **based upon the actual value of data**.

Condition Variables

Without condition variables, the programmer would need to have threads continually polling (possibly in a critical section), to check if the condition is met. This can be very resource consuming since the thread would be continuously busy in this activity. A condition variable is a way to achieve the same goal without polling.

A condition variable is always used in conjunction with a mutex lock.

Condition Variables

Thread A

- Do work up to the point where a certain condition must occur (such as "count" must reach a specified value)
- Lock associated mutex and check value of a global variable
- Call **pthread_cond_wait()** to perform a blocking wait for signal from Thread-B. Note that a call to **pthread_cond_wait()** automatically and atomically unlocks the associated mutex variable so that it can be used by Thread-B.
- When signalled, wake up. Mutex is automatically and atomically locked.
- Explicitly unlock mutex
- Continue

Thread B

- Do work
- Lock associated mutex
- Change the value of the global variable that Thread-A is waiting upon.
- Check value of the global Thread-A wait variable. If it fulfills the desired condition, signal Thread-A.
- Unlock mutex.
- Continue

Condition Variables

Condition variables must be declared with type **pthread_cond_t**, and must be initialized before they can be used. There are two ways to initialize a condition variable:

1. Statically, when it is declared. For example:
pthread_cond_t myconvar = PTHREAD_COND_INITIALIZER;
2. Dynamically, with the **pthread_cond_init()** routine.

Waiting and Signaling on Condition Variables

`pthread_cond_wait()` blocks the calling thread until the specified *condition* is signalled.

This routine should be called while *mutex* is locked, and it will automatically release the mutex while it waits.

After signal is received and thread is awakened, *mutex* will be automatically locked for use by the thread.

The programmer is then responsible for unlocking *mutex* when the thread is finished with it.

Example

condition_var.c

Example

Condition Variables and Prime numbers

`find_prime_number_seq_0.c`

Exercise 4c

Create parallel version of program for finding prime number under certain position. Use at least **two threads**.

template: find_prime_number_seq_4.c

Completing other threads

```
#include <pthread.h>
```

```
int pthread_cancel (pthread_t thread);
```

0 - if successful, otherwise the parameter **thread** is incorrect

Cancellation policy

The cancellation status of the thread (enabled by default)

```
#include <pthread.h>
```

```
int pthread_setcancelstate (int state, int *oldstate);
```

state: **PTHREAD_CANCEL_ENABLE** или
PTHREAD_CANCEL_DISABLE

Type about the thread

- Asynchronous:

At any time

- Deferred:

Only at certain moments

Change the type of cancellation

```
#include <pthread.h>
```

```
int pthread_setcanceltype (int type, int *oldtype);
```

type: **PTHREAD_CANCEL_ASYNCHRONOUS** или
PTHREAD_CANCEL_DEFERRED

Check point for delayed thread cancellation

```
#include <pthread.h>
```

```
void pthread_testcancel(void);
```

Release resources before completion

`pthread_cleanup_push(void *func, void *arg)`

Run when:

- Thread cancel
- Exit via `pthread_exit`
- `pthread_cleanup_pop`

Example

Delayed thread stop

pthread2.c

Example streams from Java

```
1 Thread t;  
2  
3 t.start();  
4 t.interrupt() // isInterrupted  
5 t.join();  
~
```

Special method `t.interrupted()`

```
if (Thread.interrupted()) {  
    throw new InterruptedException();  
}
```

Release resources

```
1 while( true ) {  
2     try {  
3         // do something  
4         // sleep(1)  
5     }  
6     catch( InterruptedException e ) {  
7  
8     }  
9 }
```

Analogue `pthread_cleanup_push`