

# OpenMP

Open Multi-Processing

# OpenMP (Open Multi-Processing)

Open standard for parallelization of programs in C, C++, Fortran languages

Interface standard for parallel programming on shared memory

Brief history:

- Version 1.0 1997
- Version 4.0 July 2013
- Version 4.1 in progress

Compilers: GCC, Intel C++ compiler, Visual C++

# Advantages Of OpenMP

- Ease of use
- Cross-platform for shared memory systems
- Hiding low-level operations
- Support for parallel and serial versions of programs

# Core components

- Environment variable
- Compiler directive
- Functions


# Steps of development

1. Write and debug a sequential program
2. Add OpenMP directives to the program
3. Use compiler with support for OpenMP
4. Set the environment variables
5. Run the program

# Compilation

```
gcc source.c -o source.x -fopenmp
```

# Model Fork-Join



- Explicit indication of parallel sections
- Support for nested parallelism
- Dynamic threads support

# Simple program

```
1  #include <omp.h>
2  #include <stdio.h>
3
4  void main() {
5      #pragma omp parallel
6      printf("Hellow word\n");
7  }
```



# Environment variable

**OMP\_NUM\_THREADS** - Sets the number of threads in the parallel block. By default, the number of threads is equal to the number of virtual processors.

**OMP\_DYNAMIC** - Allows or disables dynamic changes in the number of threads that are actually used for calculations (depending on system load). The default value depends on the implementation.

**OMP\_NESTED** - Allows or disables nested parallelism (parallelization of nested loops). The default is disabled.

# Functions of OpenMP

- **omp\_set\_num\_threads** To set the number of threads
- **omp\_get\_num\_threads** Return the number of threads in the group
- **omp\_get\_max\_threads** Maximum number of threads
- **omp\_get\_thread\_num** ID of thread
- **omp\_get\_num\_procs** Maximum number of processors
- **omp\_in\_parallel** Check for location in a parallel region

# Directives (Directive Format (C / C ++))

**#pragma omp** < name> [ <attributes> {[,] <attributes>}]

name - the name of the **Directive**;

Attributes - a construction specifying additional information and depending on the **Directive**;

# Directive parallel

**#pragma omp parallel [<attributes>]**

**<structural block>**

- **A thread that encounters a parallel construct creates a thread group, becoming the master.**
- **Threads are assigned unique integer numbers starting with 0 (the main thread).**
- **Each thread executes code defined by a building block, at the end of which a barrier is implicitly set**

# Declaring a parallel section

```
#include <omp.h>
int main()
{
    // Sequential code
    #pragma omp parallel
    {
        // Parallel code
    }
    // Sequential code

    return 0;
}
```

# Условное объявление параллельной секции

```
#include <omp.h>
int main()
{
    // Sequential code
    #pragma omp parallel if (expr)
    {
        // Parallel code
    }
    // Sequential code

    return 0;
}
```

# The parallel Directive. Simple example

```
1  #include <omp.h>
2  #include <stdio.h>
3
4  void main() {
5      #pragma omp parallel
6      printf("Hellow word\n");
7  }
```

# Examples

omp\_hello.c


omp\_hello\_functions.c

omp\_hello\_env.c




# Ways to divide work between threads


```
#pragma omp for
for (i=0;i<N;i++)
{
    // code
}
```



```
#pragma omp sections
{
    #pragma omp section
    // code 1
    #pragma omp section
    // code 2
}
```



```
#pragma omp single
{
    // code
}
```



# Directive **for**

**#pragma omp for [<attributes>]**

**<loop body>**

**Restrictions:**

- The only counter is an integer, pointer, or random access iterator. Should only change in the loop header.
- Loop condition: comparing a variable with a loop-invariant expression using <, <=, >, >=.
- Changing the counter: using ++,--, i += d, i -= d, i = i + d, i = d + i, i = i - d (d — cycle-invariant integer expression)

# Directive for

```
#include <stdio.h>
#include <omp.h>
int main()
{ int i;
  #pragma omp parallel
  {
    #pragma omp for
    for (i=0;i<1000;i++)
      printf("%d ",i);
  }
  return 0;
}
```

# Directive **for**

```
#include <stdio.h>
#include <omp.h>
int main()
{ int i;

    #pragma omp parallel for
        for (i=0;i<1000;i++)
            printf("%d ",i);

    return 0;
}
```

# Directive sections

```
#include <stdio.h>
#include <omp.h>
int main()
{ int i;
  #pragma omp parallel sections private(i)
  {
    #pragma omp section
    { printf("1st half\n");
      for (i=0;i<500;i++) printf("%d ",i);
    }
    #pragma omp section
    { printf("2nd half\n");
      for (i=501;i<1000;i++) printf("%d ",i);
    }
  }
  return 0;
}
```

In this case, you need to keep in mind the need to support the serial version of the program.

# Directive omp single

```
#include <stdio.h>
#include <omp.h>
int main()
{ int i;
  #pragma omp parallel private(i)
  {
    #pragma omp for
    for (i=0;i<1000;i++) printf("%d ",i);
    #pragma omp single
    printf("I'm thread %d!\n",get_thread_num());
    #pragma omp for
    for (i=0;i<1000;i++) printf("%d ",i);
  }
  return 0;
}
```

# Directive omp single

```
#include <stdio.h>
#include <omp.h>
int main()
{ int i;
  #pragma omp parallel private(i)
  {
    #pragma omp for
    for (i=0;i<1000;i++) printf("%d ",i);
    #pragma omp single
    printf("I'm thread %d!\n",get_thread_num());
    #pragma omp for
    for (i=0;i<1000;i++) printf("%d ",i);
  }
  return 0;
}
```

barrier

# Directive omp single

```
#include <stdio.h>
#include <omp.h>
int main()
{ int i;
  #pragma omp parallel private(i)
  {
    #pragma omp for
    for (i=0;i<1000;i++) printf("%d ",i);
    #pragma omp single nowait
    printf("I'm thread %d!\n",get_thread_num());
    #pragma omp for
    for (i=0;i<1000;i++) printf("%d ",i);
  }
  return 0;
}
```

No barrier



# Directive omp master

```
#include <stdio.h>
#include <omp.h>
int main()
{ int i;
  #pragma omp parallel private(i)
  {
    #pragma omp for
    for (i=0;i<1000;i++) printf("%d ",i);
    #pragma omp master
    printf("I'm Master!\n");
    .....
    #pragma omp for
    for (i=0;i<1000;i++) printf("%d ",i);
  }
  return 0;
}
```

No barrier

# The scope of variables

Variables declared inside a parallel block are local to the thread:

```
#pragma omp parallel
{
    int num = omp_get_thread_num();
    printf("Thread %d\n", num);
}
```

# The scope of variables

Variables declared outside the parallel block **are shared by default** across all threads:

```
int n = 10;
#pragma omp parallel
{
    for (int i=0;i<n;i++)
        printf("%d\n",i);
}
```

```
int num;
#pragma omp parallel
{
    num = omp_get_thread_num();
    printf("Thread %d\n",num);
}
```

# The scope of variables

The scope of variables declared outside the parallel block is determined by the Directive parameters:

- private
- firstprivate
- lastprivate
- shared
- default
- reduction
- threadprivate
- copyin

# The scope of variables

The scope of variables declared outside the parallel block is determined by the Directive parameters:

- **private** - Own local variable in each thread
- firstprivate
- lastprivate
- shared
- reduction

```
int num;  
#pragma omp parallel private(num)  
{  
    num = omp_get_thread_num();  
    printf("%d\n", num);  
}
```

# The scope of variables

The scope of variables declared outside the parallel block is determined by the Directive parameters:

- private
- **firstprivate** - Local variable with initialization
- lastprivate
- shared
- reduction

```
int num = 5;
#pragma omp parallel \
    firstprivate(num)
{
    printf("%d\n", num);
}
```

# The scope of variables

The scope of variables declared outside the parallel block is determined by the Directive parameters:

- private
- firstprivate
- **lastprivate**
- shared
- reduction

Local variable with the last value saved (in sequential execution)

```
int i,j;  
#pragma omp parallel for \  
                    lastprivate(j)  
for (i=0;i<100;i++) j = i;  
  
printf("Последний j = %d\n",j);
```

# The scope of variables

The scope of variables declared outside the parallel block is determined by the Directive parameters:

- private
- firstprivate
- lastprivate
- **shared**
- reduction

Shared (shared) variable

```
int i,j;  
  
#pragma omp parallel for \  
    shared(j)  
for (i=0;i<100;i++) j = i;  
  
printf("j = %d\n",j);
```



# The scope of variables

The scope of variables declared outside the parallel block is determined by the Directive parameters:

- private
- firstprivate
- lastprivate
- shared
- **reduction**

The variable to perform the reducing operation

```
int i,s = 0;
#pragma omp parallel for \
        reduction(+:s)
    for (i=0;i<100;i++)
        s += i;

printf("Sum: %d\n",s);
```

# Synchronizing threads

## Thread synchronization directives:

- master
- barrier
- critical
- atomic

## Locks:

- `omp_lock_t`

# Synchronizing threads

## Thread synchronization directives:

- **master**
- barrier
- critical
- atomic

Выполнение кода только главным потоком

```
#pragma omp parallel
{
    //code
    #pragma omp master
    {
        // critical code
    }
    // code
}
```

# Synchronizing threads

## Thread synchronization directives:

- master
- **barrier**
- critical
- atomic

```
#pragma omp parallel
{
    printf("Hello!\n");

    #pragma omp barrier
    printf("I am thread %d\n",
          omp_get_thread_num());
}
```

# Synchronizing threads

## Thread synchronization directives:

- master
- barrier
- **critical**
- atomic

```
int i,idx[N],x[M];

#pragma omp parallel for
for (i=0;i<N;i++)
{
    #pragma omp critical
    {
        x[idx[i]] += count(i);
    }
}
```

# Synchronizing threads

## Thread synchronization directives:

- master
- barrier
- critical
- **atomic**

```
int i,idx[N],x[M];

#pragma omp parallel for
for (i=0;i<N;i++)
{
    #pragma omp atomic
    x[idx[i]] += count(i);
}
```

# Locks

- Lock-a special object common to threads
- Threads can capture (lock) and release (unlock) a lock
- Only one thread at a time can capture a lock
- When attempting to capture a lock, threads wait for the lock to be released
- Using locks, you can control access to shared resources

# An example of using the lock

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
int x[1000];
int main()
{ int i,max;
  omp_lock_t lock;
  omp_init_lock(&lock);
  for (i=0;i<1000;i++) x[i]=rand();
  max = x[0];
  #pragma omp parallel for
    for(i=0;i<1000;i++)
    { omp_set_lock(&lock);
      if (x[i]>max) max = x[i];
      omp_set_unlock(&lock);
    }
  omp_destroy_lock(&lock);
  return 0;
}
```



# Example: single vs critical

```
1  int a=0, b=0;
2  #pragma omp parallel num_threads(4)
3  {
4      #pragma omp single
5      a++;
6      #pragma omp critical;
7      b++;
8  }
9
10 printf("single: %d -- critical: %d\n", a, b);
```

# Example: single vs critical

## Output:


single: 1 -- critical: 4

# Example

## Numerical integration (Riemann sum)

Numerical integration is the approximate computation of an integral using numerical techniques

Source: omp\_integral\_4.c



# Exercise 5a

Create program for Pi calculation based on OpenMP. Implement two versions:

1. Without reduction
2. Use reduction attribute

Questions:

Compare execution time of two versions.